

Yongho Nam (Mark)

Mississauga, ON • yongho.nam@mail.utoronto.ca • github.com/dydggh2011 • linkedin.com/in/yongho-nam

EXPERIENCE

Gamespring

Mar 2025 – Aug 2026

Backend Developer, Admin Platforms & AI

South Korea

- Built an AI agent pipeline for development, QA, deployment, and review after the team shrank from three engineers to one, shipping **5 features a week** solo for 18 months, up from 3 for the former team, with only review and release done by hand.
- Built multi-agent orchestration with delegated sub-agents before native tool support, running **3–4 tasks in parallel**.
- Developed canary checks and rule-splitting for long-running agents, flagging when an agent began dropping its constraints so the session could be compressed or restarted before faulty work reached review.
- Piloted AWS Transform on a 20-feature admin page with AWS solutions staff, benchmarking it against Claude Code on coverage, accuracy, usability, and cost; presented the case against it to the CEO after costs ran **10x over estimate**.
- Proposed and presented to the CEO a self-improving agent that read the internal project tracker to model engineer capability and auto-route work, winning approval to continue the research.
- Ran the admin platform itself as sole developer, including 6 months on the AI team: 300K lines and 150 features across revenue analytics, live control of 8 game systems, player accounts, and events, used daily by **90% of staff**.
- Shipped **40 new features, 50 upgrades, and 200+ bug fixes** while maintaining 90% of the platform, learning React from scratch to ship a first front-end feature in 3 days.
- Built a self-serve console that let 6 client engineers switch branches across 180 dev server processes themselves, removing an 8-minute wait on backend staff from 12 switches a day (**8 engineer-hours a week**).
- Interviewed the customer service team on how they searched players, then moved pagination ahead of multi-table joins and built them a dedicated search view, cutting filtered searches from 60 s to **10 s**, and to **under 1 s** on the new view.
- Fixed a stamp-challenge reset that locked 8 game tables for 2 minutes, freezing or disconnecting live players; scoped it to affected tables and reset players lazily on login, cutting it to **under 2 seconds**.
- Triaged 4 requests a week from across the company, holding scoping meetings with requesters when build cost outweighed value; trained new hires on the platform and handed it to the next owner over 3 months with full documentation.

Gamespring

Jul 2024 – Mar 2025

Backend Developer, Game Server

South Korea

- Shipped seasonal mode Christmas Jackpot in **1 day instead of 3** by reusing the payout engine behind a new routing layer.
- Built a time-series heap monitor that found the cause of a production-only memory leak in 2 days; the leak grew a live slot server by 3 GB an hour and forced 2 restarts a week, and **none recurred** after the fix.

PROJECTS

Kelly | *TypeScript, Electron, Node.js, SQLite, Ollama, OpenAI API*

Code/docs on request • Sept 2026 – Present

- Building an open-source, local-first macOS app that logs active apps and keyboard/mouse activity for an AI (local or the user's API key) to score against their focus goals.
- Keeps activity data on device for users to audit, and enforces missed goals with penalties such as automatic stock purchases.

The Nerdiest Dungeon | *Godot, TypeScript, Node.js, PostgreSQL, AWS*

Code/docs on request • Aug 2026 – Present

- Designing a browser-based pixel-art roguelike RPG that turns studying into turn-based combat: learning goals become dungeon roadmaps, monsters are question sets, and correct answers fire skills scaled by difficulty.
- Specified gamified spaced repetition, where earned items expire unless related review questions are solved, and a JSON content layer for community-authored learning paths with AI-generated fallbacks.

What To Cook | *TypeScript, NestJS, PostgreSQL, TypeORM, OpenAI API, AWS EC2, GitHub Actions* Mar – May 2024

- Built a full-stack recipe platform combining the OpenAI API with Google Custom Search, with request throttling, HTML sanitization, and DTO validation, turning a user's on-hand ingredients into dish suggestions and sourced recipes.
- Ran it as a live service behind a GitHub Actions pipeline to AWS EC2 with Swagger docs and Jest/Supertest end-to-end tests, routing runtime errors to Slack through Sentry before users could report them.

EDUCATION

University of Toronto

Honours Bachelor of Science, Computer Science and Mathematics

Toronto, ON

Expected 2029

SKILLS

AI Engineering: Multi-agent orchestration, sub-agent delegation, context engineering, agent rule adherence, prompt design, AI-assisted development, AWS Bedrock, Claude, OpenAI API, AWS Transform, Amazon Quick Suite

Languages: TypeScript, JavaScript, SQL, Java, Go, HTML/CSS, Python **Spoken:** English, Korean

Backend: Node.js, Express, NestJS, Socket.IO, REST, WebSocket, TypeORM, Passport/JWT/OAuth, Swagger, Jest

Databases: MySQL/Aurora, PostgreSQL, Redis; EXPLAIN profiling, indexing, window functions, JSON columns, rollups

Infrastructure: AWS (EC2, RDS, S3, CloudWatch), Docker, Kubernetes, GitHub Actions, PM2, Linux, Sentry